



BITS Pilani
Pilani Campus

OS Tutorial

Interprocess communication (IPC) (shared memory)

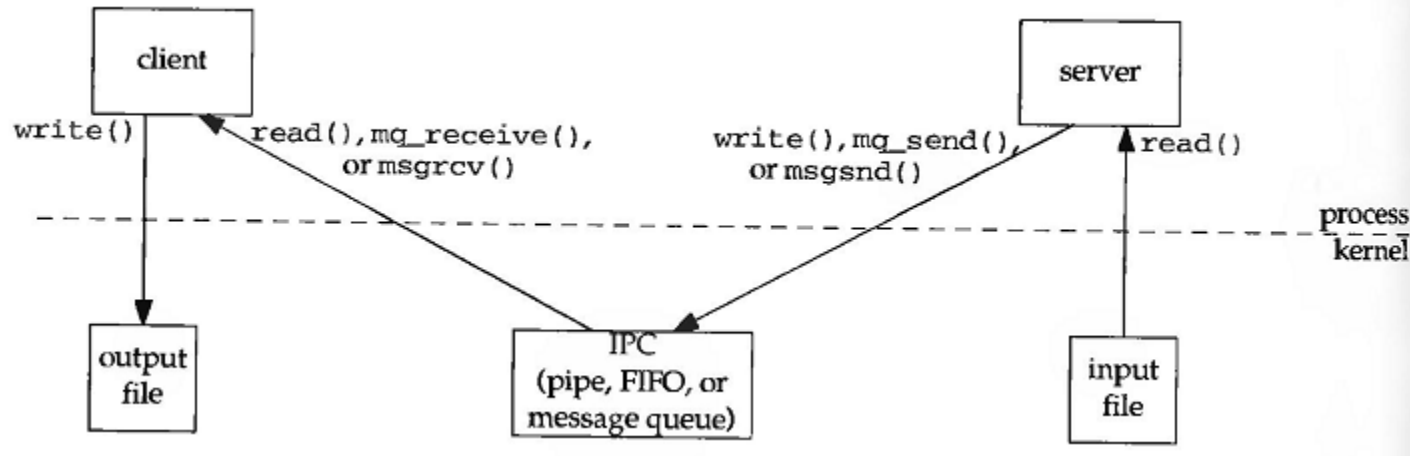
IPC methods

- Signal
- Pipe
- FIFO
- Shared Memory
- Message passing
- Semaphores

Objectives


- How the processes can communicate among themselves using the Shared Memory.
- Creating a Shared Memory Segment.
- Controlling a Shared Memory Segment.
- Attaching and Detaching a Shared Memory Segment.

Message Passing



Takes 4 copies to transfer data between two processes

Shared Memory




- ✓ *Takes only two steps*
- ✓ *Kernel is not involved in transferring data but it is involved in creating shared memory*

Shared Memory

- ✓ Shared Memory is an efficient means of passing data between programs.
- ✓ One program will create a memory portion, which other processes (if permitted) can access.
- ✓ A shared segment can be attached multiple times by the same process.
- ✓ A shared memory segment is described by a control structure with a unique ID that points to an area of physical memory.

Shared Memory



Shared Memory

- ✓ A shared memory is an extra piece of memory that is attached to some address spaces for their owners to use.
- ✓ As a result, all the processes share the same memory segment and have access to it.

Shared Memory

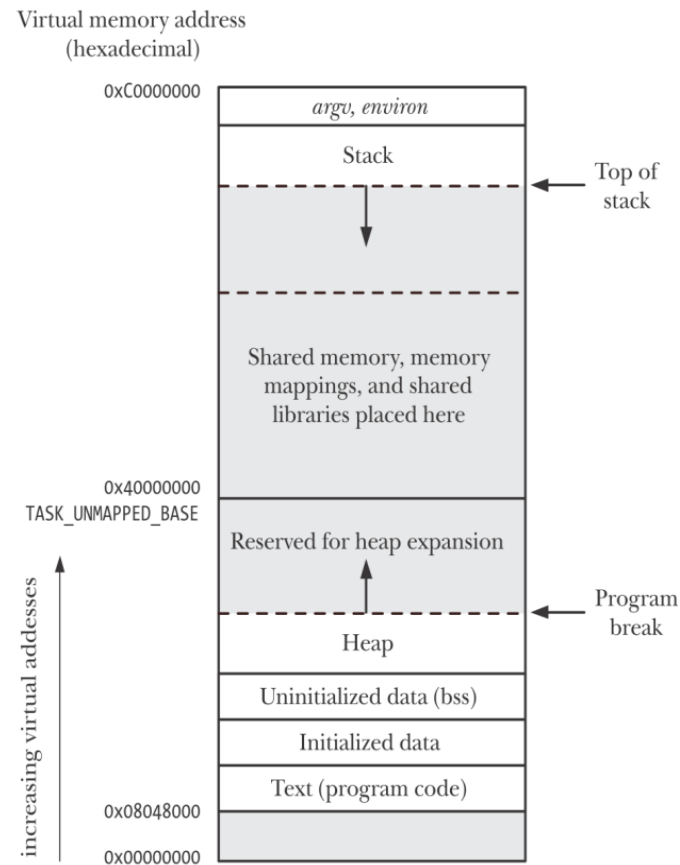
- ✓ Shared memory is a feature supported by UNIX System V, including Linux, SunOS and Solaris.
- ✓ One process must explicitly ask for an area, using a key, to be shared by other processes. This process will be called the server. All other processes, the clients that know the shared area can access it.

Shared Memory

- ✓ However, there is no protection to a shared memory and any process that knows it can access it freely.
- ✓ To protect a shared memory from being accessed at the same time by several processes, a synchronization protocol must be setup.

Location of Shared Memory in Virtual Memory

- When shared memory segment is mapped on the process address space using recommended method, it is attached as shown.



Shared Memory Limits

- SHMMNI
 - System limit on no of shared memory identifies
- SHMMIN
 - Minimum size of a shared memory segment
- SHMAX
 - Maximum size of shared memory segment.
- SHMALL
 - System limit of total number of pages of shared memory.

```

1  $ cd /proc/sys/kernel
2  $ cat shmmni
3  4096
4  $ cat shmmax
5  33554432
6  $ cat shmall
7  2097152

```

Table 48-2: System V shared memory limits

Limit	Ceiling value (x86-32)	Corresponding file in /proc/sys/kernel
SHMMNI	32768 (IPCMNI)	shmmni
SHMAX	Depends on available memory	shmmax
SHMALL	Depends on available memory	shmall

Shared Memory

- ✓ The shared memory itself is described by a **structure of type `shmid_ds`** in header file **`sys/shm.h`**.

```
1 struct shmid_ds {  
2     struct ipc_perm shm_perm;    /* Ownership and permissions */  
3     size_t    shm_segsz;        /* Size of segment in bytes */  
4     time_t    shm_atime;        /* Time of last shmat() */  
5     time_t    shm_dtime;        /* Time of last shmdt() */  
6     time_t    shm_ctime;        /* Time of last change */  
7     pid_t     shm_cpid;         /* PID of creator */  
8     pid_t     shm_lpid;         /* PID of last shmat() / shmdt() */  
9     shmatt_t  shm_nattch;       /* Number of currently attached processes */  
10 };
```

Shared Memory

- ✓ The shared memory itself is described by a **structure of type `shmid_ds`** in header file **`sys/shm.h`**.
- ✓ To use this file, following files must be included:

```
#include <sys/types.h>  
#include <sys/ipc.h>  
#include <sys/shm.h>
```

Shared Memory

- For a server, it should be started before any client.
- The server should perform the following tasks:
 - Ask for a shared memory with a memory key and memorize the returned **shared memory ID**. This is performed by **system call shmget()**.
 - Attach this shared memory to the server's address space with **system call shmat()**.
 - Initialize the shared memory, if necessary.
 - Do something and wait for all clients' completion.
 - Detach the shared memory with **system call shmdt()**.
 - Remove the shared memory with system call **shmctl()**.

Shared Memory

- For the client part, the procedure is almost the same:
 - Ask for a shared memory with the same memory key and memorize the returned shared memory ID.
 - Attach this shared memory to the client's address space.
 - Use the memory.
 - Detach all shared memory segments, if necessary.
 - Exit.

shmget()

- This system call requests a shared memory segment.
- It is defined as follows:
- **shm_id** = **shmget** (
 key_t **k**, /* the key for the segment */
 int **size**, /* the size of the segment */
 int **flag** /* create/use flag */
);
- If it is successful, it returns a non-negative integer, the shared memory ID; otherwise, the function value is negative.

shmget()

- **k** is of type **key_t** or **IPC_PRIVATE**. It is the numeric key to be assigned to the returned shared memory segment.
- **size** is the size of the requested shared memory.
- **flag** is used to specify the way that the shared memory will be used.
 - Only the following two values are important:
 - **IPC_CREAT | 0666** for a server (i.e., creating and granting read and write access to the server).
 - **0666** for any client (i.e., granting read and write access to the client)

shmget()

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
#include <stdio.h>
.....
int  shm_id; /* shared memory ID */
.....
shm_id = shmget (IPC_PRIVATE, 4*sizeof(int), IPC_CREAT | 0666);
if (shm_id < 0)
{
    printf("shmget error\n");
    exit(1);
}
```

key

- UNIX requires a key of type **key_t** defined in file **sys/types.h** for requesting shared memory segments.
- There are three different ways of using keys, namely:
 - A specific integer value (e.g., 123456)
 - A key generated with function **ftok()**
 - A uniquely generated key using **IPC_PRIVATE** (i.e., a private key).

key

- The first way is the easiest one; however, its use may be very risky since a process can access your resource as long as it uses the same key value to request that resource.
- The following example assigns 1234 to a key:
key_t SomeKey;
SomeKey = 1234;
- The **ftok()** function has the following prototype:

```
key_t ftok (  
    const char *path, /* a path string */  
    int id           /* an integer value */  
);
```

key

- `ftok()` takes a character string that identifies a path and an integer (usually a character) value, and generates an integer of type `key_t` based on the first argument with the value of `id` in the most significant position.
- Example: if the generated integer is $35028A5D_{16}$ and the value of `id` is 'a' (ASCII value = 61_{16}), then `ftok()` returns $61028A5D_{16}$. That is, 61_{16} replaces the first byte of $35028A5D_{16}$, generating $61028A5D_{16}$.

key

- ✓ If a processes use the same arguments to call `ftok()`, the returned key value will always be the same.
- ✓ The most commonly used value for the first argument is `"."`, the current directory.
- ✓ If all related processes are stored in the same directory, the following call to **`ftok()`** will generate the same key value.

```
#include <sys/types.h>
#include <sys/ipc.h>

key_t  SomeKey;
SomeKey = ftok(".", 'x');
```

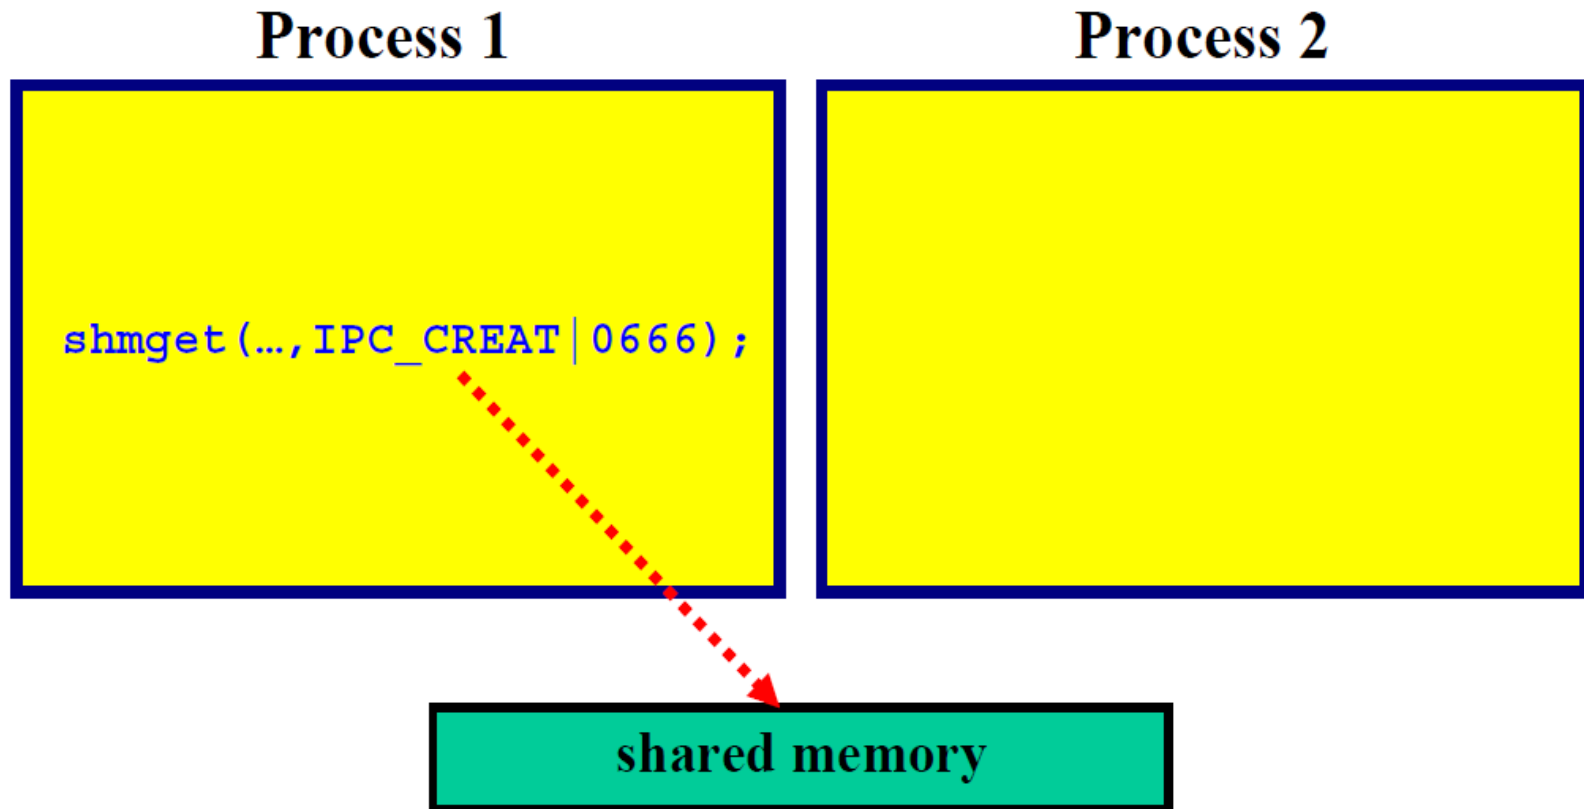

key

- After obtaining a key, it can be used in any place where a key is required.
- Moreover, the place where a key is required accepts a special parameter, IPC_PRIVATE.
 - the system will generate a unique key and guarantee that no other process will have the same key.

key

- If a resource is requested with IPC_PRIVATE in a place where a key is required, that process will receive a unique key for that resource.
 - Since that resource is identified with a unique key unknown to the outsiders, other processes will not be able to share that resource and, as a result, the requesting process is guaranteed that it owns and accesses that resource exclusively.

After the Execution of shmget()



Shared memory is allocated; but, is not part of the address space

key

- Once a shared memory segment has been created, a process attaches it to its address space by calling `shmat`.

```

1  #include <sys/types.h>          /* For portability */
2  #include <sys/shm.h>
3  void *shmat(int  shmid , const void * shmaddr , int  shmflg );
4  //Returns address at which shared memory is attached on success,
5  //or (void *) -1 on error

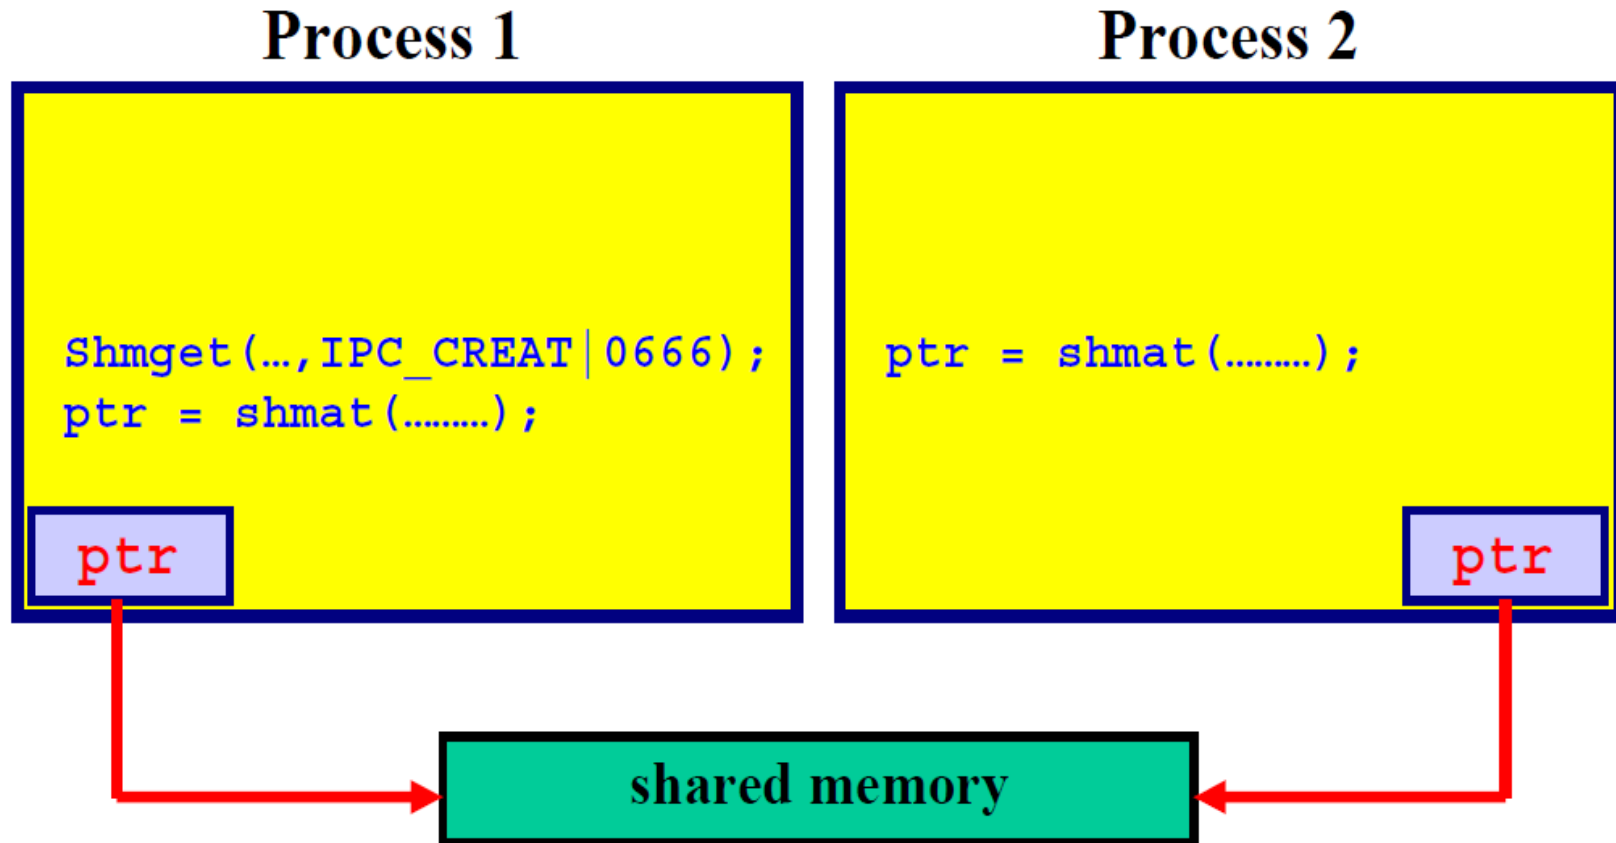
```

The address in the calling process at which the segment is attached depends on the `addr` argument.

- If `addr` is 0, the segment is attached at the first available address selected by the kernel.
- This is the recommended technique. **Table 48-1:** *shmflg* bit-mask values for *shmat()*

Value	Description
SHM_RDONLY	Attach segment read-only
SHM_REMAP	Replace any existing mapping at <i>shmaddr</i>
SHM_RND	Round <i>shmaddr</i> down to multiple of SHMLBA bytes

After the Execution of shmat()



Now processes can access the shared memory

Server Program

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
#include <stdio.h>

int shm_id;
key_t mem_key;
int *shm_ptr;
mem_key = ftok(".", 'a');
shm_id = shmget(mem_key,
    4*sizeof(int), IPC_CREAT | 0666); }

if (shm_id < 0) {
    printf("*** shmget error (server) ***\n");
    exit(1);
}

shm_ptr = (int *) shmat(shm_id, NULL, 0);
    /* attach */
if ((int) shm_ptr == -1) {
    printf("*** shmat error (server) ***\n");
    exit(1);
}
```

Client Program

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
#include <stdio.h>

int shm_id;
key_t mem_key;
int *shm_ptr;
mem_key = ftok(".", 'a');
shm_id = shmget(mem_key,
    4*sizeof(int), 0666);
```

```
if (shm_id < 0) {
    printf("*** shmget error (server) ***\n");
    exit(1);
}

shm_ptr = (int *) shmat(shm_id, NULL, 0);
    /* attach */
if ((int) shm_ptr == -1) {
    printf("*** shmat error (server) ***\n");
    exit(1);
}
```

shmdt() and shmctl()

- **shmdt() is used to detach a shared memory.**
- After a shared memory is detached, it cannot be used. However, it is still there and can be re-attached back to a process's address space.
- **To remove a shared memory, use shmctl().**
- The only argument to shmdt() is **the shared memory address returned by shmat().**

shmdt() and shmctl()

- **shmdt (shm_ptr), it will detached the shared memory.**
- shm_ptr is the pointer to the shared memory, returned by shmat() during shared memory attachment.
- **If shmdt() fails, the returned value is non-zero.**
- To remove a shared memory segment, use the following shmctl (shm_id, IPC_RMID, NULL).
- **shm_id is the shared memory ID.**
- **IPC_RMID indicates this is a remove operation.**

- Two different processes communicating via shared memory

```
main()
```

```
{
```

```
char c;
```

```
int shmid;
```

```
key_t key;
```

```
char *shm, *s;
```

```
key = 5678;
```

```
/* * create the segment. * */
```

```
if ((shmid = shmget(key, SHMSIZE, IPC_CREAT | 0666)) < 0)
```

```
{
```

```
perror("shmget");    exit(1);
```

```
}
```

```
/** Now we attach the segment to our data space. */
```

```
if ((shm = shmat(shmid, NULL, 0)) == (char *) -1)
```

```
{
```

```
perror("shmat");
```

```
exit(1);
```

```
}
```

Server Process

```
/** Now put some things into the memory for the other process to read. */  
s = shm;  
for (c = 'a'; c <= 'z'; c++)  
*s++ = c;  
*s = NULL;  
/** Finally, we wait until the other process  
* Changes the first character of our memory  
* to '*', indicating that it has read what  
* we put there.  
*/  
while (*shm != '*')  
sleep(1);  
exit(0);  
}
```

Server Process

```
main()
{
int shmid;

key_t key;
char *shm, *s;
/*
 * We need to get the segment named
 * "5678", created by the server.
 */
key = 5678;
/*
 * Locate the segment.
 */
if ((shmid = shmget(key, SHMSIZE, 0666)) < 0) {
perror("shmget");
exit(1);
}
```

Client Process

```
/* Now we attach the segment to our data space*/  
if ((shm = shmat(shmid, NULL, 0)) == (char *) -1)  
{  
perror("shmat");  
exit(1);  
}  
/* Now read what the server put in the memory*/  
for (s = shm; *s != NULL; s++)  
    putchar(*s);  
    putchar('\n');  
/* Finally, change the first character of the  
* segment to '*', indicating we have read  
* the segment.  
*/  
*shm = '*';  
printf ("\nIts done from client.\n\n\n");  
exit(0);  
}
```

Client Process

- Parent and Child processes communicating via shared memory



```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
```

```
int main(void)
{
    int shmid;
    char *shmPtr;
    int n;
    if (fork( ) == 0)
    {
        sleep(5);
        if( (shmid = shmget(2041, 32, 0)) == -1 )
        {
            exit(1);
        }
    }
}
```



```
shmPtr = shmat(shmid, 0, 0);
if (shmPtr == (char *) -1)
    exit(2);
printf ("\nChild Reading ....\n\n");
for (n = 0; n < 26; n++)
    putchar(shmPtr[n]);
    putchar('\n');
}
else
{
    if( (shmid = shmget(2041, 32, 0666 | IPC_CREAT)) == -1 )
    {
        exit(1);
    }
    shmPtr = shmat(shmid, 0, 0);
```

```
if (shmPtr == (char *) -1)
    exit(2);
for (n = 0; n < 26; n++)
    shmPtr[n] = 'a' + n;
printf ("Parent Writing ....\n\n");
for (n = 0; n < 26; n++)
    putchar(shmPtr[n]);
    putchar('\n');
wait(NULL);
shmdt(NULL);
if( shmctl(shmid, IPC_RMID, NULL) == -1 )
{
    perror("shmctl");
    exit(-1);
}
exit(0);
}
```

```
if (shmPtr == (char *) -1)
    exit(2);
for (n = 0; n < 26; n++)
    shmPtr[n] = 'a' + n;
printf ("Parent Writing ....\n\n");
for (n = 0; n < 26; n++)
    putchar(shmPtr[n]);
    putchar('\n');
wait(NULL);
shmdt(NULL);
if( shmctl(shmid, IPC_RMID, NULL) == -1 )
{
    perror("shmctl");
    exit(-1);
}
exit(0);
}
```